



Optimization and Regularization of Neural Networks

Tim Ransom



College of
**ENGINEERING, COMPUTING
AND APPLIED SCIENCES**





Outline

Optimization

Regularization of NN





Outline

Optimization

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization

Regularization of NN

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- Dropout



Learning vs. Optimization

Goal of learning: minimize generalization error, or the loss function

$$\mathcal{L}(W) = \mathbb{E}_{(x,y) \sim p_{data}} \left[L(f(x, W), y) \right]$$

In practice, empirical risk minimization:


$$\mathcal{L}(W) = \sum_i [L(f(x_i; W), y_i)]$$

f is the neural network

Quantity optimized
different from the quantity
we care about



Local Minima





Critical Points

Points with zero gradient

2nd-derivate (Hessian) determines curvature




Image Source: Goodfellow et al. (2016)




Local Minima

Old view: local minima is major problem in neural network training

Recent view:

- For sufficiently large neural networks, **most local minima incur low cost**
- Not important to find true global minimum



Saddle Points

Recent studies indicate that in high dim, saddle points are more likely than local min

Gradient can be very small near saddle points




Image Source: Goodfellow et al. (2016)




Poor Conditioning

Poorly conditioned Hessian matrix

- **High curvature:** small steps leads to huge increase

Learning is slow despite strong gradients

Oscillations slow down progress






No Critical Points

Some cost functions do not have critical points. In particular: classification.

WHY?





Exploding and Vanishing Gradients




Image Source: deeplearning.ai

Linear activation $h_i = Wx$
 $h_i = Wh_{i-1}, \quad i = 2, \dots, n$



Exploding and Vanishing Gradients

Suppose $\mathbf{W} = \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$:

$$\begin{bmatrix} h_1^1 \\ h_2^1 \end{bmatrix} = \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \quad \dots \quad \begin{bmatrix} h_1^n \\ h_2^n \end{bmatrix} = \begin{bmatrix} a^n & 0 \\ 0 & b^n \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$



Exploding and Vanishing Gradients

Suppose $x = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$

Case 1: $a = 1, b = 2$:

$$y \rightarrow 1, \quad \nabla y \rightarrow \begin{bmatrix} n \\ n2^{n-1} \end{bmatrix} \quad \text{Explodes!}$$

Case 2: $a = 0.5, b = 0.9$:

$$y \rightarrow 0, \quad \nabla y \rightarrow \begin{bmatrix} 0 \\ 0 \end{bmatrix} \quad \text{Vanishes!}$$



Exploding and Vanishing Gradients

Exploding gradients lead to cliffs

Can be mitigated using [gradient clipping](#)

if $\|g\| > u$

$$g \leftarrow \frac{gu}{\|g\|}$$




Image Source: Goodfellow et al. (2016)



Outline

Optimization

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization


Regularization of NN

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- Dropout



Momentum

Oscillations because updates do not exploit curvature information




Average gradient presents faster path to optimal:
vertical components cancel out



Momentum


Question: Why not this?





Momentum


Let us figure out an algorithm which will lead us to the minimum faster.





Momentum


Look each component at a time





Momentum


Let us figure out an algorithm





Momentum


Let us figure out an algorithm





Momentum


Let us figure out an algorithm





Momentum

Let us figure out an algorithm





Momentum

f is the Neural Network

Old gradient descent:

$$g = \frac{1}{m} \sum_i \nabla_W L(f(x_i; W), y_i)$$

$$W^* = W - \lambda g$$

$$g = g + \text{average } g \text{ from before}$$

New gradient descent with momentum:

$$v = \alpha v + (1 - \alpha) g$$

$$W^* = W - \lambda v$$

$\alpha \in [0, 1)$ controls how quickly
effect of past gradients decay



Nesterov Momentum

Apply an **interim** update:

$$\tilde{W} = W + \nu$$

Perform a correction based on gradient at the interim point:

$$g = \frac{1}{m} \sum_i \nabla_W L(f(x_i; \tilde{W}), y_i)$$

$$\nu = \alpha \nu - \varepsilon g$$

$$W = W + \nu$$

Momentum based on
look-ahead slope



Outline

Optimization

- Challenges in Optimization
- Momentum
- **Adaptive Learning Rate**
- Parameter Initialization
- Batch Normalization


Regularization of NN

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- Dropout





Adaptive Learning Rates




Oscillations along vertical direction

- Learning must be slower along parameter 2

Use a different learning rate for each parameter?



Adaptive Learning Rates




Oscillations along vertical direction

- Learning must be slower along parameter 2

Use a different learning rate for each parameter?



Adaptive Learning Rates




Oscillations along vertical direction

- Learning must be slower along parameter 2

Use a different learning rate for each parameter?



Adaptive Learning Rates




Oscillations along vertical direction

- Learning must be slower along parameter 2

Use a different learning rate for each parameter?



Adaptive Learning Rates



Oscillations along vertical direction

- Learning must be slower along parameter 2

Use a different learning rate for each parameter?



AdaGrad

- Accumulate squared gradients:

$$r_i = r_i + g_i^2$$

g is the gradient

- Update each parameter:

$$W_i = W_i - \frac{\epsilon}{\delta + \sqrt{r_i}} g_i$$

Inversely proportional to
cumulative squared gradient

- Greater progress along gently sloped directions



AdaGrad

δ is a small number, making sure this does not become too large

Old gradient descent:

$$g = \frac{1}{m} \sum_i \nabla_W L(f(x_i; W), y_i)$$

W^* $W - \lambda g$

We would like λ 's not to be the same and inversely proportional to the $|g_i|$

$$W_i^* = W_i - \lambda_i g_i \qquad \lambda_i \propto \frac{1}{|g_i|} = \frac{1}{\delta + |g_i|}$$

New gradient descent with adaptive learning rate:

$$r_i^* = r_i + g_i^2 \qquad W_i^* = W_i - \frac{\epsilon}{\delta + \sqrt{r_i}} g_i$$





RMSProp

- For non-convex problems, AdaGrad can prematurely decrease learning rate
- Use **exponentially weighted average** for gradient accumulation

$$r_i = \rho r_i + (1 - \rho) g_i^2$$

$$W_i = W_i - \frac{\epsilon}{\delta + \sqrt{r_i}} g_i$$



Adam

- RMSProp + Momentum
- Estimate first moment:

$$v_i = \rho_1 v_i + (1 - \rho_1) g_i$$

- Estimate second moment:

$$r_i = \rho_2 r_i + (1 - \rho_2) g_i^2$$

- Update parameters:

$$W_i = W_i - \frac{\epsilon}{\delta + \sqrt{r_i}} v_i$$

Works well in practice,
is fairly robust to
hyper-parameters

Also applies
bias correction
to v and r



Outline

Optimization

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- **Parameter Initialization**
- Batch Normalization

Regularization of NN

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- Dropout



Parameter Initialization

- Goal: **break symmetry** between units
 - so that each unit computes a different function
- Initialize all weights (not biases) **randomly**
 - Gaussian or uniform distribution
- **Scale of initialization?**
 - *Large* -> grad explosion, *Small* -> grad vanishing



Xavier Initialization

- Heuristic for all outputs to have **unit variance**
- For a fully-connected layer with m inputs:

$$W_{ij} \sim N\left(0, \frac{1}{m}\right)$$

- For ReLU units, it is recommended:

$$W_{ij} \sim N\left(0, \frac{2}{m}\right)$$



Normalized Initialization

- Fully-connected layer with m inputs, n outputs:

$$W_{ij} \sim U\left(-\sqrt{\frac{6}{m+n}}, \sqrt{\frac{6}{m+n}}\right)$$

- Heuristic trades off between initialize all layers have same activation and gradient variance
- **Sparse** variant when m is large
 - Initialize k nonzero weights in each unit



Outline

Optimization

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization

Regularization of NN

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- Dropout



Feature Normalization

Good practice to normalize features before applying learning algorithm:

$$x' = \frac{x - \mu}{\sigma}$$

Feature vector $\rightarrow$ x Vector of mean feature values $\rightarrow$ μ
Vector of SD of feature values $\rightarrow$ σ


Features in **same scale**: mean 0 and variance 1

- Speeds up learning




Feature Normalization

$L(W)$



Before normalization




After normalization



Internal Covariance Shift

Each hidden layer changes distribution of inputs to next layer: *slows down learning*



Normalize
inputs to layer 2

...

Normalize
inputs to layer n



Batch Normalization

Training time:

- Mini-batch of activations for layer to normalize

$$H = \begin{bmatrix} \boxed{H_{11} \cdots H_{1K}} \\ \vdots \quad \ddots \quad \vdots \\ \boxed{H_{N1} \cdots H_{NK}} \end{bmatrix} \begin{matrix} K \text{ hidden layer} \\ \text{activations} \end{matrix}$$

N data points
in mini-batch



Batch Normalization

Training time:

- Mini-batch of activations for layer to normalize

$$H' = \frac{H - \mu}{\sigma}$$

where

$$\mu = \frac{1}{m} \sum_i H_{i,:}$$

Vector of mean activations
across mini-batch

$$\sigma = \sqrt{\frac{1}{m} \sum_i (H - \mu)_i^2 + \delta}$$

Vector of SD of each unit
across mini-batch



Batch Normalization

Training time:

- Normalization can reduce expressive power
- Instead use:

$$\gamma H' + \beta$$


Learnable parameters

- Allows network to **control range of normalization**




Batch Normalization

Batch 1



$$\mu^1 = \frac{1}{m} \sum_i H_{i,:}$$
$$\sigma^1 = \sqrt{\frac{1}{m} \sum_i (H - \mu)_i^2 + \delta}$$

Batch N




Add normalization
operations for
layer 1




Batch Normalization

Batch 1



$$\mu^2 = \frac{1}{m} \sum_i H_{i,:}$$
$$\sigma^2 = \sqrt{\frac{1}{m} \sum_i (H - \mu)_i^2 + \delta}$$

Batch N



Add normalization
operations for
layer 2 and so on ...



Outline

Optimization

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization

Regularization of NN

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- Dropout



Regularization

Regularization is any modification we make to a learning algorithm that is intended to **reduce its generalization error** but not its training error.



Outline

Optimization

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization


Regularization of NN

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- Dropout



Overfitting

Fitting a deep neural network with 5 layers and 100 neurons per layer can lead to a very good prediction on the training set but poor prediction on validation set.





Norm Penalties

We used to optimize:

$$L(W; X, y)$$

Change to ...

$$L_R(W; X, y) = L(W; X, y) + \alpha \Omega(W)$$

Biases not
penalized

L_2 regularization:

- Weights decay
- MAP estimation with Gaussian prior

L_1 regularization:

- encourages sparsity
- MAP estimation with Laplacian prior

$$\Omega(W) = \frac{1}{2} \| W \|_2^2$$

$$\Omega(W) = \frac{1}{2} \| W \|_1$$



Norm Penalties

We used to optimize:

$$W^{(i+1)} = W^{(i)} - \lambda \frac{\partial L}{\partial W}$$

Change to ...

$$L_R(W; X, y) = L(W; X, y) + \frac{1}{2} \alpha W^2$$

$$W^{(i+1)} = W^{(i)} - \lambda \frac{\partial L}{\partial W} - \lambda \alpha W$$

Weights
decay in
proportion to
size

Biases not
penalized

L_2 regularization:

- Decay of weights
- MAP estimation with Gaussian prior

L_1 regularization:


- encourages sparsity
- MAP estimation with Laplacian prior

$$\Omega(W) = \frac{1}{2} \|W\|_2^2$$


$$\Omega(W) = \frac{1}{2} \|W\|_1$$



Norm Penalties



$$\Omega(W) = \frac{1}{2} \|W\|_2^2$$



$$\Omega(W) = \frac{1}{2} \|W\|_1$$



Outline

Optimization

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization


Regularization of NN

- Norm Penalties
- Early Stopping**
- Data Augmentation
- Sparse Representation
- Dropout



Early Stopping


Early stopping: terminate while validation set performance is better



Training time can be treated as a hyperparameter



Early Stopping





Outline

Optimization

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization


Regularization of NN

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- Dropout





Data Augmentation





Outline

Optimization


- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization

Regularization of NN

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- Dropout



Sparse Representation




$$L(\theta; X, y)$$

$$[4.34] = \underbrace{[3.2 \quad 2.0 \quad 1.8]}_{W_7} \begin{bmatrix} 2 \\ -2.2 \\ 1.3 \end{bmatrix}$$



Sparse Representation




$$L(\theta; X, y)$$

$$[4.34] = [3.2 \quad 2 \quad 1] \left[\begin{array}{c} 2 \\ -2.2 \\ 1.3 \end{array} \right] \left. \vphantom{\begin{array}{c} 2 \\ -2.2 \\ 1.3 \end{array}} \right\} h_{31}, h_{32}, h_{33}$$



Sparse Representation




$$L_R(W; X, y) = L(\theta; X, y) + \alpha \Omega(W)$$

$$[0.69] = \underbrace{[0.5 \quad .2 \quad 0.1]}_{W_7} \begin{bmatrix} 2 \\ -2.2 \\ 1.3 \end{bmatrix}$$

Weights in output layer



Sparse Representation



$$L_R(W; X, y) = L(\theta; X, y) + \alpha \Omega(h)$$

$$[1.3] = [3.2 \quad 2 \quad 1] \begin{bmatrix} 0 \\ -0.2 \\ .9 \end{bmatrix} \quad \left. \vphantom{\begin{bmatrix} 0 \\ -0.2 \\ .9 \end{bmatrix}} \right\} h_{31}, h_{32}, h_{33}$$

Output of hidden layer



Outline

Optimization

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization

Regularization of NN

- Norm Penalties
- Early Stopping
- Data Augmentation
- Sparse Representation
- Dropout



Noise Robustness

Random **perturbation of network weights**

- Gaussian noise: Equivalent to minimizing loss with regularization term
- Encourages smooth function: small perturbation in weights leads to small changes in output


Injecting **noise in output labels**

- Better convergence: prevents pursuit of hard probabilities




Dropout

- Randomly set some neurons and their connections to zero (i.e. “dropped”)
- Prevent overfitting by reducing co-adaptation of neurons
- Like training many random sub-networks



(a) Standard Neural Net




(b) After applying dropout.



Dropout

- Widely used and highly effective
- Proposed as an alternative to ensembling, which is too expensive for neural nets



Test error for different architectures with and without dropout. The networks have 2 to 4 hidden layers each with 1024 to 2048 units.

Image Source:

<http://jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf>



Dropout

- Widely used and highly effective
- Proposed as an alternative to ensembling, which is too expensive for neural nets

Test error for different architectures with and without dropout. The networks have 2 to 4 hidden layers each with 1024 to 2048 units.





Image Source:

<http://jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf>





Questions?



School of
COMPUTING
Clemson University